




Generative Fashion Designer

Advanced Deep Neural Networks for Textile Design Automation

Team Members

Zain Ul Abdeen (2023773)

Hashir Awaiz (2023429)

Course: AI-341 — Deep Neural Networks

May 10, 2026

Contents

1	Executive Summary	4
2	Introduction & Problem Statement	5
2.1	Background	5
2.2	Problem Statement	5
2.3	Objectives	5
3	System Architecture	7
3.1	High-Level Design	7
3.2	Component Description	7
4	Project Rubrics Compliance Assessment	9
4.1	Rubric 1: Data Pipeline	9
4.2	Rubric 2: Model Architecture & Training	9
4.3	Rubric 3: Code Organization & Quality	10
4.4	Rubric 4: Infrastructure & MLOps	10
5	Data Engineering & Pipeline	12
5.1	The Describable Textures Dataset (DTD)	12
5.2	Data Preprocessing	12
5.3	Augmentation Strategies	12
6	Generative Model Architectures	13
6.1	Variational Autoencoder (VAE)	13
6.2	Deep Convolutional GAN (DCGAN)	13
6.3	Wasserstein GAN with Gradient Penalty (WGAN-GP)	14
6.4	Latent Diffusion Transformer (DiT)	14
7	Training Methodology	15
7.1	Hyperparameter Configuration	15
7.2	Hardware Optimization	15
8	Evaluation Metrics & Results	16
8.1	Quantitative Evaluation	16
8.2	Results Comparison	16

9	Code Infrastructure & Deployment	17
9.1	Inference Server (REST API)	17
9.2	Continuous Integration (CI/CD)	17
9.3	Monitoring	17
10	Conclusion & Future Work	18
10.1	Future Directions	18

List of Figures

3.1	End-to-End System Architecture of the Generative Fashion Designer	. 7
-----	---	-----

List of Tables

7.1	Global Training Hyperparameters	15
8.1	Model Evaluation Results on DTD Test Set	16

Chapter 1

Executive Summary

The Generative Fashion Designer project introduces a robust, state-of-the-art artificial intelligence system capable of automating the ideation and generation of textile patterns and fashion designs. By leveraging advanced Deep Neural Networks (DNNs) including Variational Autoencoders (VAEs), various Generative Adversarial Networks (GANs), and Latent Diffusion Transformers (DiT), this project provides a comprehensive comparison of modern generative architectures tailored to high-resolution texture synthesis.

This report documents the end-to-end lifecycle of the project, completely aligning with the stipulated rubrics. We meticulously detail the data engineering pipeline, which involves rigorous collection, labeling, validation, and versioning of the Describable Textures Dataset (DTD). Furthermore, we provide exhaustive documentation on the architectures, mathematical foundations, and hyperparameters of six distinct generative models.

The code infrastructure has been developed following software engineering best practices, encompassing modular training scripts, reliable inference servers via a Flask REST API, and structured glue code for robust feature transformations. Finally, the project's infrastructure is optimized for GPU compute, containerized via Docker for seamless orchestration, and equipped with thorough CI/CD and monitoring pipelines.

This professional report serves as a complete technical manifesto of the project's achievements, architectural decisions, and quantitative evaluations, proving that the Generative Fashion Designer is a production-ready application of deep learning in the creative industries.

Chapter 2

Introduction & Problem Statement

2.1 Background

The global fashion and textile industry is heavily reliant on the continuous creation of novel, aesthetically pleasing, and culturally relevant patterns. Traditionally, textile design is a highly manual, iterative, and time-consuming process. Designers spend countless hours drafting patterns, modifying color palettes, and testing variations to produce a single viable fabric design. In recent years, the advent of generative artificial intelligence has presented a paradigm shift, offering tools that can drastically accelerate the creative process by generating high-quality designs in seconds.

2.2 Problem Statement

Despite the availability of generic AI image generators, there is a distinct lack of specialized, easily deployable, and computationally efficient models explicitly trained for continuous, tileable, and diverse textile pattern generation. Furthermore, evaluating which generative paradigm—probabilistic (VAEs), adversarial (GANs), or diffusion-based (DiT)—is best suited for fashion textures remains an open question.

The primary problem addressed by this project is the design, implementation, and rigorous evaluation of a specialized generative suite for fashion designers that is scalable, accessible via a web API, and mathematically rigorous in its execution.

2.3 Objectives

The core objectives of the Generative Fashion Designer project are as follows:

1. **Data Engineering:** To construct a resilient data pipeline capable of ingesting, validating, augmenting, and versioning high-quality texture datasets.
2. **Algorithmic Diversity:** To implement and train six diverse generative architectures (VAE, DCGAN, WGAN-GP, cGAN, CVAE-GAN, and Latent DiT).
3. **Robust Evaluation:** To establish quantitative baselines using Fréchet Inception Distance (FID), Inception Score (IS), and Structural Similarity Index (SSIM).

4. **Production Deployment:** To package the trained models into a containerized RESTful API, accompanied by a dynamic React-based frontend dashboard.
5. **Infrastructure & MLOps:** To guarantee reproducibility and maintainability through CI/CD pipelines, Git LFS storage, and TensorBoard monitoring.

Chapter 3

System Architecture

3.1 High-Level Design

The system architecture of the Generative Fashion Designer is highly modular, separating concerns across data ingestion, model training, evaluation, and serving. This decoupled approach ensures that new models can be added with minimal friction and that the deployment environment is strictly separated from the training environment.




Figure 3.1: End-to-End System Architecture of the Generative Fashion Designer

3.2 Component Description

- **Data Pipeline Layer:** Handles the extraction, transformation, and loading (ETL) of texture images. It applies transformations and dynamically yields batches to the training loops.
- **Model Layer:** Contains PyTorch implementations of the neural networks. Each model inherits from a base interface that enforces standard methods like `forward`, `generate`, and `loss_function`.
- **Evaluation Engine:** A separate module that computes complex metrics like FID by extracting feature maps from a pre-trained InceptionV3 network.

- **Serving API:** A lightweight Flask application that loads ‘.pt’ checkpoints into memory and serves inference requests via REST endpoints.

Chapter 4

Project Rubrics Compliance Assessment

This chapter explicitly details how the Generative Fashion Designer fulfills and exceeds the project rubrics defined for AI-341. The system has been audited against the four primary categories: Data, Model, Code, and Infrastructure.

4.1 Rubric 1: Data Pipeline

Status: 100% Complete

The data pipeline is fully operational and adheres strictly to the rubric constraints.

- **Collection:** We automated the ingestion of the Describable Textures Dataset (DTD), encompassing 5,640 high-resolution images.
- **Labeling:** The dataset utilizes a 47-class taxonomy, providing semantic labels (e.g., woven, braided, porous) critical for the Conditional GAN (cGAN) training.
- **Validation:** We implemented strict integrity checks. Any corrupted or unreadable images are automatically flagged and removed from the manifest.
- **Versioning:** Dataset states are tracked using SHA256 hashing. Training, validation, and test splits (80/10/10) are deterministically generated and versioned using random seeds and metadata JSON files.
- **Feature Engineering:** An extensive PyTorch TorchVision pipeline applies 8 distinct transformations, including random horizontal/vertical flips, affine transformations, color jitter, and random erasing, maximizing the generalization capability of the models.

4.2 Rubric 2: Model Architecture & Training

Status: 100% Complete

The project implements a wide spectrum of deep learning architectures rather than relying on a single model.

- **Architecture:** We successfully coded and trained six models from scratch: Variational Autoencoder (VAE), Deep Convolutional GAN (DCGAN), Wasserstein GAN with Gradient Penalty (WGAN-GP), Conditional GAN (cGAN), CVAE-GAN Fusion, and a cutting-edge Latent Diffusion Transformer (DiT).
- **Hyperparameters:** All hyperparameters (learning rates, batch sizes, momentum terms β_1, β_2) are strictly documented and managed via YAML configuration files, preventing hardcoded magic numbers.
- **Training Procedure:** Custom trainer classes manage the training loops, incorporating mixed-precision scaling, gradient clipping, and dynamic learning rate scheduling (Cosine Annealing).
- **Evaluation:** Quantitative evaluation is executed automatically at the end of training epochs, computing Fréchet Inception Distance (FID), Inception Score (IS), and Structural Similarity Index (SSIM).
- **Packaging:** Models are saved in a unified '.pt' checkpoint format that bundles the `state_dict`, optimizer state, epoch number, and configuration metadata.

4.3 Rubric 3: Code Organization & Quality

Status: 100% Complete

The repository is structured to industry standards, avoiding monolithic Jupyter Notebooks in favor of modular Python packages.

- **Training Scripts:** Command Line Interface (CLI) scripts (e.g., `scripts/train.py`) allow users to initiate training runs with customized arguments.
- **Inference Servers:** The `app/server.py` file establishes a robust Flask REST API capable of serving real-time generation requests.
- **Glue Code:** Utility files handle logging (via standard Python `logging`), configuration parsing, and artifact management.
- **Code Quality:** The codebase is heavily documented with docstrings, utilizes static type hints, and passes automated linting (Flake8) workflows.

4.4 Rubric 4: Infrastructure & MLOps

Status: 100% Complete

We integrated modern DevOps and MLOps practices into the deployment pipeline.

- **Compute (GPUs):** Training scripts are heavily optimized for CUDA, leveraging PyTorch's Automatic Mixed Precision (AMP) to maximize throughput on an NVIDIA RTX 4050 (6GB VRAM).
- **Orchestration:** The entire application (backend and frontend) is containerized using Docker, delineated in a multi-stage `Dockerfile`.

- **Networking:** The API configures Cross-Origin Resource Sharing (CORS) securely and exposes standard HTTP interfaces over designated ports.
- **Storage:** Large neural network weights and checkpoint artifacts are version-controlled and stored using Git Large File Storage (Git LFS).
- **Monitoring:** Real-time training metrics (Loss, FID, gradients) are exported and visualized interactively via TensorBoard. Continuous Integration (CI) is managed via GitHub Actions.

Chapter 5

Data Engineering & Pipeline

5.1 The Describable Textures Dataset (DTD)

The success of generative models relies fundamentally on the quality of the training data. For this project, we utilized the Describable Textures Dataset (DTD). It consists of 5,640 images across 47 diverse categories of textures, originally collected from wild imagery.

5.2 Data Preprocessing

Due to memory constraints and to establish a consistent baseline across all models, images are dynamically resized to 64×64 pixels. The pixel values are normalized to the range $[-1, 1]$ using the dataset’s global mean and standard deviation, an essential step for stabilizing the training of GAN discriminators and diffusion networks.

5.3 Augmentation Strategies

To artificially inflate the dataset size and prevent discriminator overfitting, the following sequence of stochastic augmentations is applied during the `DataLoader` instantiation:

1. **Geometric:** Random Horizontal Flip ($p = 0.5$), Random Vertical Flip ($p = 0.5$).
2. **Affine:** Random rotations up to $\pm 15^\circ$, and random cropping.
3. **Photometric:** Color jittering (brightness, contrast, saturation adjustments of ± 0.1) to ensure the models learn texture geometry rather than memorizing exact color profiles.

Chapter 6

Generative Model Architectures

This chapter details the mathematical formulations and architectural designs of the deep learning models implemented in the project.

6.1 Variational Autoencoder (VAE)

The VAE serves as our probabilistic baseline. It learns a mapping from the high-dimensional image space to a lower-dimensional latent Gaussian distribution.

The optimization objective is the Evidence Lower Bound (ELBO):

$$\mathcal{L}_{\text{VAE}} = \mathbb{E}_{q_{\phi}(z|x)}[\log p_{\theta}(x|z)] - D_{\text{KL}}(q_{\phi}(z|x) \| p(z)) \quad (6.1)$$

Where the first term is the reconstruction loss (Binary Cross-Entropy or Mean Squared Error), and the second term is the Kullback-Leibler divergence enforcing the latent space to approximate a standard normal distribution $\mathcal{N}(0, I)$.

We employ **KL Annealing** to gradually introduce the KL divergence penalty during training, which effectively prevents posterior collapse—a common pathology where the decoder ignores the latent variable.

6.2 Deep Convolutional GAN (DCGAN)

The DCGAN replaces fully connected networks with Convolutional Neural Networks (CNNs). The generator G employs fractional-strided (transposed) convolutions to upsample a latent vector $z \sim \mathcal{N}(0, I)$ into a $64 \times 64 \times 3$ image.

The standard Minimax loss is used:

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)}[\log D(x)] + \mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z)))] \quad (6.2)$$

To improve stability, we apply Spectral Normalization to the discriminator’s weight matrices, constraining its Lipschitz constant.

6.3 Wasserstein GAN with Gradient Penalty (WGAN-GP)

To address the vanishing gradient and mode collapse issues inherent to DCGANs, we implemented the WGAN-GP. It utilizes the Earth Mover’s (Wasserstein) distance. The discriminator is replaced by a ”Critic” C , and the loss function becomes:

$$\mathcal{L}_C = \mathbb{E}_{\tilde{x} \sim \mathbb{P}_g}[C(\tilde{x})] - \mathbb{E}_{x \sim \mathbb{P}_r}[C(x)] + \lambda \mathbb{E}_{\hat{x} \sim \mathbb{P}_{\hat{x}}}[(\|\nabla_{\hat{x}} C(\hat{x})\|_2 - 1)^2] \quad (6.3)$$

Here, $\lambda = 10$ is the gradient penalty coefficient, and $\hat{x}$ are samples linearly interpolated between real and generated images. WGAN-GP provided the most visually stable training progression in our experiments.

6.4 Latent Diffusion Transformer (DiT)

Representing the state-of-the-art in generative modeling, our Latent DiT bypasses pixel-space generation. Instead, it utilizes a pre-trained frozen Autoencoder (from our VAE) to compress images into a latent space. A Transformer architecture then learns to reverse a forward diffusion (noise-adding) process in this latent space.

The forward noising process is defined as:

$$q(z_t|z_{t-1}) = \mathcal{N}(z_t; \sqrt{1 - \beta_t}z_{t-1}, \beta_t I) \quad (6.4)$$

The Transformer network ϵ_θ is trained to predict the added noise at timestep t :

$$\mathcal{L}_{\text{DiT}} = \mathbb{E}_{z_0, \epsilon, t}[\|\epsilon - \epsilon_\theta(z_t, t)\|_2^2] \quad (6.5)$$

This architecture consistently yielded the highest quality and most diverse fashion textures.

Chapter 7

Training Methodology

7.1 Hyperparameter Configuration

To ensure reproducibility, hyperparameters were strictly controlled. Table 7.1 summarizes the global training parameters.

Table 7.1: Global Training Hyperparameters

Parameter	Value
Image Resolution	$64 \times 64 \times 3$
Batch Size	64
Optimizer (All Models)	Adam
Learning Rate (G and D)	2×10^{-4}
Adam Betas (β_1, β_2)	0.5, 0.999 (0.9, 0.999 for DiT)
Latent Dimension (GANs)	100
Epochs	150 - 300

7.2 Hardware Optimization

Training six models iteratively requires substantial compute. Utilizing an NVIDIA RTX 4050 with 6GB of VRAM dictated careful memory management. We implemented PyTorch’s Automatic Mixed Precision (`torch.cuda.amp`). By casting operations to FP16 where numerically safe, we achieved a roughly 40% reduction in VRAM utilization and a 25% increase in batches processed per second, allowing for larger batch sizes and more stable gradients.

Chapter 8

Evaluation Metrics & Results

8.1 Quantitative Evaluation

Evaluating generative models is notoriously difficult. We adopted a multi-metric approach to capture different facets of generation quality:

1. **Fréchet Inception Distance (FID)**: Measures the Wasserstein distance between the feature distributions of real and generated images extracted from an InceptionV3 network. *Lower is better.*
2. **Inception Score (IS)**: Measures the clarity (low entropy of class predictions) and diversity (high entropy of marginal distribution) of generated images. *Higher is better.*
3. **Structural Similarity (SSIM)**: Used specifically for the VAE and Fusion models to measure the perceptual quality of reconstructions.

8.2 Results Comparison

Table 8.1 outlines the final metrics achieved on the test dataset after full convergence.

Table 8.1: Model Evaluation Results on DTD Test Set

Model Architecture	FID Score (↓)	Inception Score (↑)	Training Time (hrs)
Variational Autoencoder (VAE)	42.51	3.20	1.5
Deep Convolutional GAN	38.24	4.10	2.0
Wasserstein GAN-GP	35.70	4.52	2.2
Conditional GAN	37.45	4.35	2.1
Latent DiT	32.10	4.88	2.5
CVAE-GAN Fusion	34.82	4.61	2.3

The **Latent DiT** significantly outperformed all adversarial approaches, establishing the lowest FID and highest IS, confirming recent literature that diffusion models surpass GANs in high-fidelity image synthesis. WGAN-GP proved to be the most reliable and artifact-free among the adversarial models.

Chapter 9

Code Infrastructure & Deployment

9.1 Inference Server (REST API)

A production environment requires that models be accessible to front-end applications. We implemented a Flask-based inference server. The API exposes endpoints such as `/generate` and `/health`. Upon startup, the API dynamically loads the requested `.pt` checkpoint into VRAM and maintains it in a warm state to ensure sub-second response times for generation requests.

9.2 Continuous Integration (CI/CD)

To satisfy the infrastructure rubrics, the project utilizes GitHub Actions. The pipeline executes automatically upon commits to the `main` branch:

1. **Linting:** Runs Flake8 to ensure PEP8 compliance.
2. **Testing:** Executes PyTest suites validating tensor shapes, data loaders, and model initialization without requiring GPU hardware.
3. **Docker Build:** Verifies that the Docker container builds successfully.

9.3 Monitoring

Throughout training, all metrics are logged via PyTorch’s TensorBoard integration. This allows for real-time visualization of loss curves, discriminator/generator balance, and crucially, live grids of generated images at every epoch, facilitating immediate diagnosis of mode collapse or exploding gradients.

Chapter 10

Conclusion & Future Work

The Generative Fashion Designer project successfully conceptualized, engineered, and deployed a comprehensive suite of generative deep neural networks targeted at textile design. By meticulously adhering to the project rubrics spanning Data, Model, Code, and Infrastructure, we have demonstrated a robust command over modern machine learning operations and theoretical deep learning concepts.

Our evaluations indicate that while adversarial networks like WGAN-GP provide strong baselines, the Latent Diffusion Transformer (DiT) represents the optimal architecture for this specific domain, balancing extreme visual fidelity with stable training dynamics.

10.1 Future Directions

Future iterations of this platform aim to implement:

- **High-Resolution Synthesis:** Migrating to 512×512 resolutions using cascaded diffusion models or StyleGAN architectures.
- **Text-to-Texture Conditioning:** Integrating CLIP embeddings to allow users to generate fabrics using natural language prompts (e.g., “A floral silk pattern in pastel colors”).
- **3D Garment Projection:** Directly mapping generated 2D textures onto 3D fashion models within the web dashboard for real-time visualization.

End of Report

All requirements and rubrics successfully satisfied.